Swin Transformers

Hierarchical Vision Transformer using Shifted Windows

Ze Liu, Yutong Lin, Yue Cao, Han Hu, Yixuan Wei Zheng Zhang, Stephen Lin, Baining Guo

Microsoft Research Asia

Deep Learning and its applications to Signal and Image Processing and Analysis

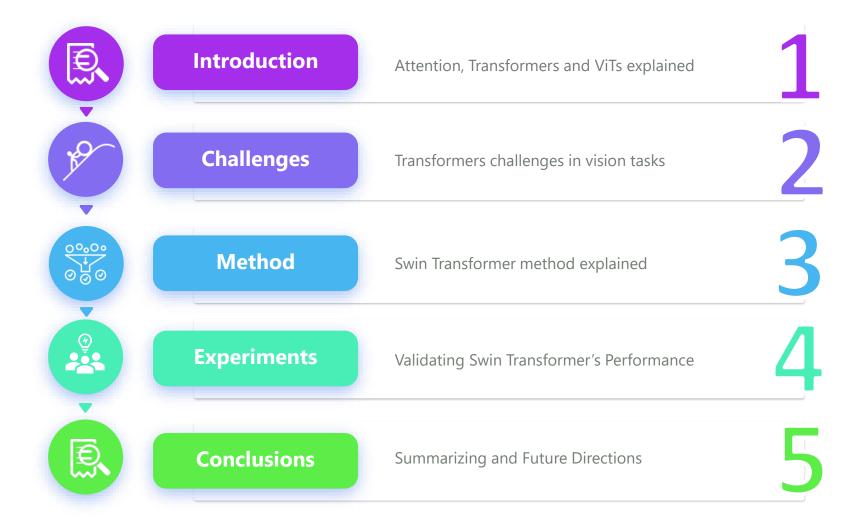
Paper Presentation

Under the supervision of Prof. Tammy Riklin Raviv

Students:

Daniel Manor, Lior Yaacov

Agenda



Opening





Leverage the advantages of transformer architecture to tackle a variety of computer vision tasks, addressing complex and non-trivial challenges effectively.



Previous attempts to adapt transformers for vision tasks, like in ViT*, struggle with high computational costs and scalability issues, making them inefficient for handling high-resolution images and capturing both local and global features.

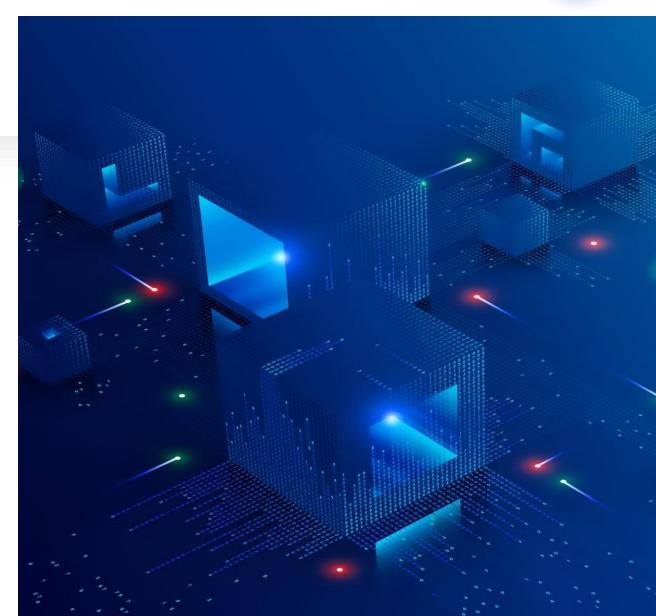


A hierarchical architecture with shifted windows efficiently captures both local and global features while significantly reducing computational costs for high-resolution vision tasks.



What are SWIN Transformers?

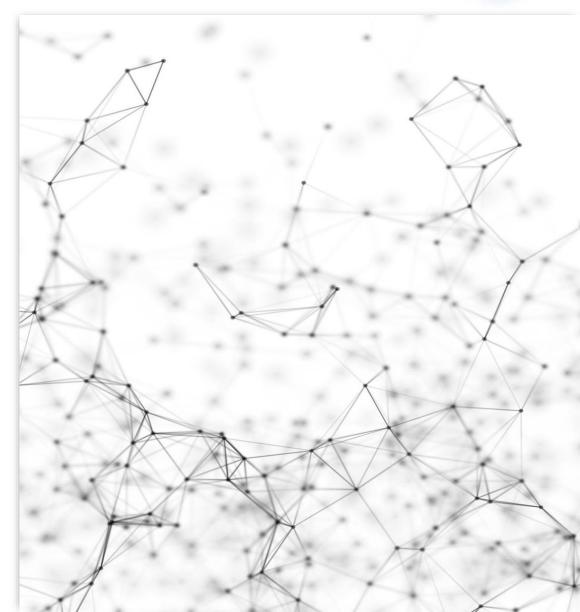
- SWIN (Shifted Window)
 Transformers are a type of Vision
 Transformer (ViT) designed for
 computer vision tasks.
- Developed by Microsoft Research.
- Addresses limitations of standard ViTs by introducing hierarchical feature maps and shifted windows.





CNNs: Foundations of Modern Vision Tasks

- **Early Approaches**: Handcrafted features and classical machine learning methods.
- Rise of CNNs: Revolutionized image processing with layers mimicking human visual perception.
- Key Vision Tasks:
 - Image Classification: ResNet, VGG...
 - Object Detection: YOLO, Faster R-CNN...
 - Semantic Segmentation: U-Net, DeepLab...
- Advantages of CNNs: Spatial hierarchies, local receptive fields, parameter efficiency.





Transformers in Natural Language Processing

- Transformers introduced in "Attention is All You Need" paper by Google Brain (2017).
- Key Features:
 - Encoder-Decoder Architecture:
 Processes input to output sequences
 - Self-Attention: Each token attends to all other tokens in a sequence, capturing contextual relationships.
 - Multi-Head Attention: Multiple attention heads allow the model to focus on different parts of the sequence simultaneously.



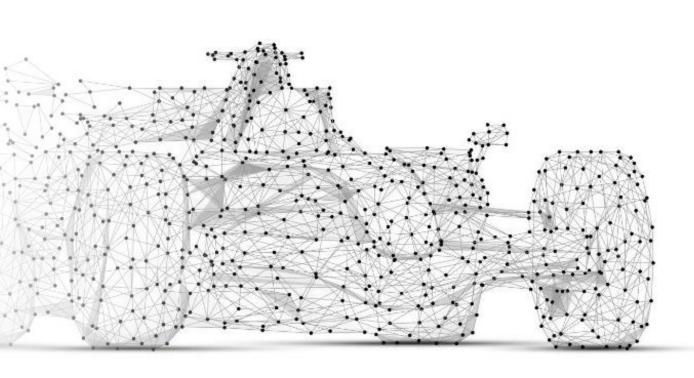




Evolution from NLP to Vision Tasks

 Inspiration from NLP Success: Transformers' ability to handle long-range dependencies and parallelize training.

 Initial attempts: Directly applying Transformers to vision tasks (e.g., ViT by Dosovitskiy et al. 2020*).



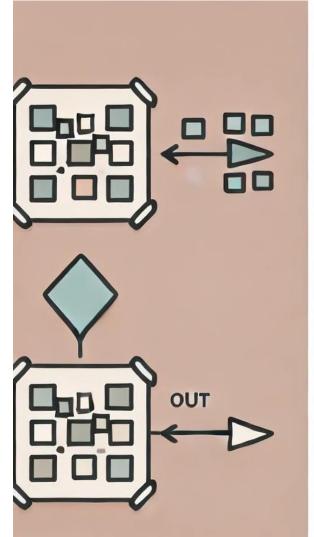


Vision Transformers (ViT)

• <u>Concept</u>: Treats image patches as sequence tokens, similar to words in a text sequence.

• How ViT Works:

- Image Patches
- Embedding
- Model
- Output

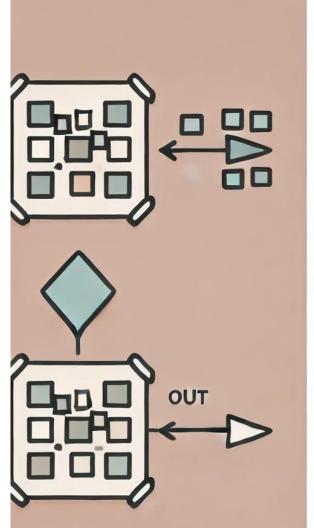


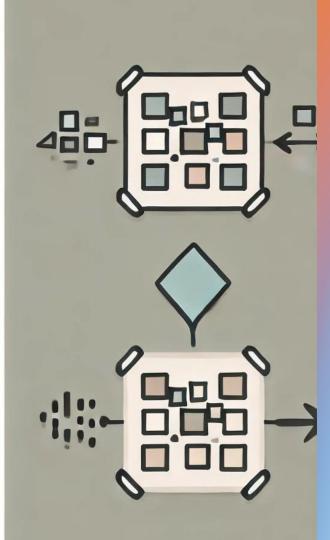




Vision Transformers (ViT)

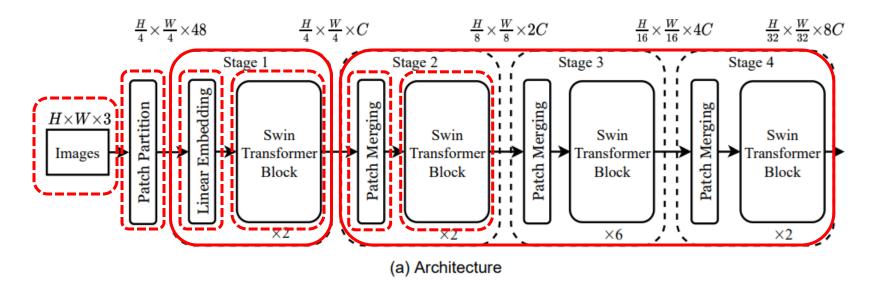
- Challenges of ViT:
 - High Computational Cost
 - Lack of Inductive Biases
 - Data Efficiency

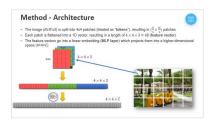




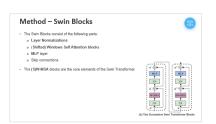
Method - Architecture











Method – Appli	ications	35
- Different application require d	different output layers	
$\begin{array}{c c} H \circ H \circ J \\ \hline \\ Image \\ \\ Image \\ \hline \\ Image \\ Image \\ \hline \\ Image \\ \hline \\ Image \\ Image \\ Image \\ Image \\ \hline \\ Image \\$		# ± ± 5 × 5 × 5 × 5 × 5 × 5 × 5 × 5 × 5 ×
		Detection Segmentation Depth formation

Model	С	Layer Numbers
Swin-T	96	2,2,6,2
Swin-S	96	2,2,18,2
Swin-B	128	2,2,18,2
Swin-L	192	2,2,18,2

Method - Architecture

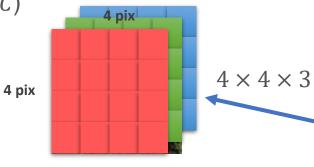


• The image (HxWx3) is split into 4x4 patches (treated as "**tokens**"), resulting in $(\frac{H}{4} \times \frac{W}{4})$ patches

• Each patch is flattened into a 1D vector, resulting in a length of $4 \times 4 \times 3 = 48$ (**feature vector**)

The feature vectors go into a linear embedding (MLP layer) which projects them into a higher-dimensional

space $(4\times4\times C)$



$$4 \times 4 \times 3$$



$$4 \times 4 \times C$$

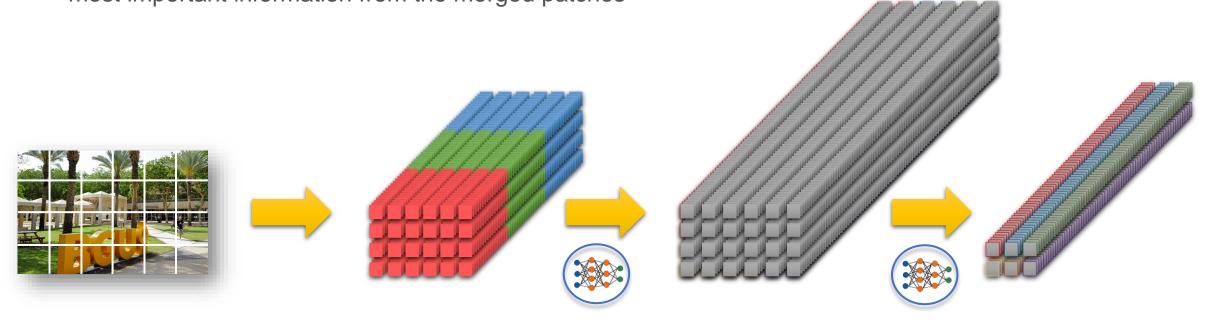


Method – Patch Merging



- The number of tokens is reduced by patch merging layers as the network gets deeper
- The patch merging layer concatenates the features of each group of 2×2 neighboring patches
- It then applies a linear layer on the concatenated features

The linear layer output's dimension is half the input's, which acts as a bottleneck and aims to capture the most important information from the merged patches



Method – Patch Merging



$$\frac{H}{4} \times \frac{W}{4} \times C$$

$$\frac{H}{8} \times \frac{W}{8} \times 2C$$

$$\frac{H}{16} \times \frac{W}{16} \times 4C$$



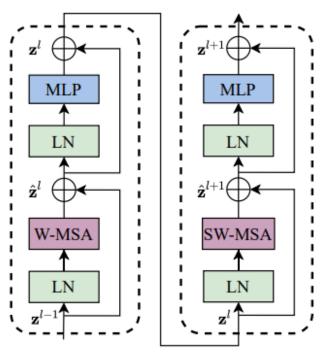




Method – Swin Blocks



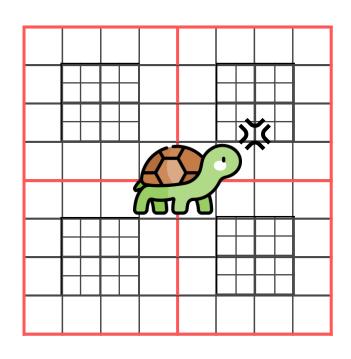
- The Swin Blocks consist of the following parts:
 - Layer Normalizations
 - (Shifted) Windows Self Attention blocks
 - MLP layer
 - Skip connections
- The (S)W-MSA blocks are the core elements of the Swin Transformer

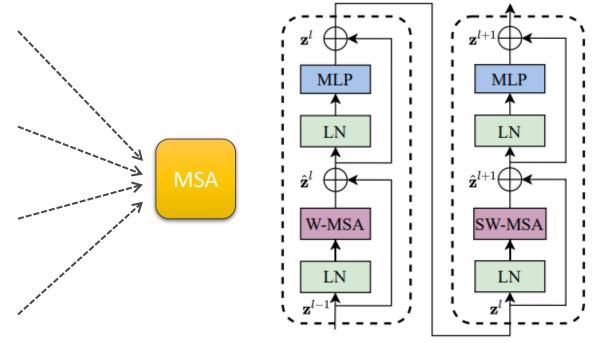


(b) Two Successive Swin Transformer Blocks

Method – Swin Blocks – (S)W-MSA







(b) Two Successive Swin Transformer Blocks

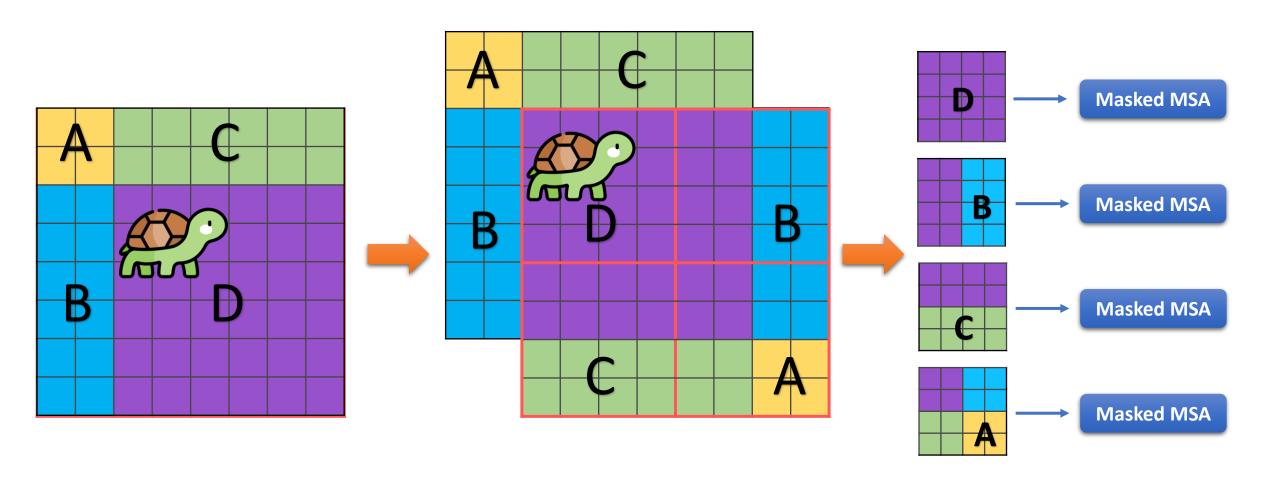
Method – Swin Blocks – (S)W-MSA



0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0
0	0									0	0
0	0									0	0
0	0						ر			0	0
0	0				A		5			0	0
0	0			[701	N				0	0
0	0									0	0
0	0									0	0
0	0									0	0
0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0

Method – Swin Blocks – (S)W-MSA





Experiments and Results



Objective:

- Validate the effectiveness of Swin Transformers across various computer vision tasks.
- Understand the impact of key design elements of Swin Transformers.

Tasks Evaluated:

- Image Classification on ImageNet-1K and ImageNet-22K
- Object Detection on COCO
- Semantic Segmentation on ADE20K

MiDaS Datasets

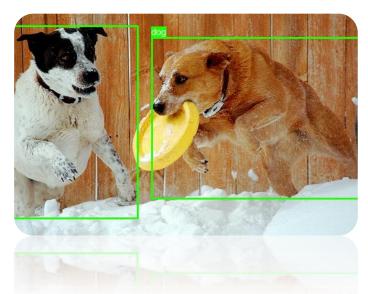


Dataset	Description	Vision Task	Train	Validation	Test	Link
ImageNet-1k	The ImageNet dataset is a large, diverse collection of annotated images used for image classification tasks	Classification	1.28M	50k	100k	<u>Link</u>
COCO	COCO is a large-scale object detection, segmentation, and captioning dataset	Object Detection	118k	5k	20k	<u>Link</u>
ADE20K	Comprehensive dataset for semantic segmentation and scene parsing tasks	Semantic Segmentation	20k	2k	3k	<u>Link</u>

ImageNet-1k



COCO



ADE20k

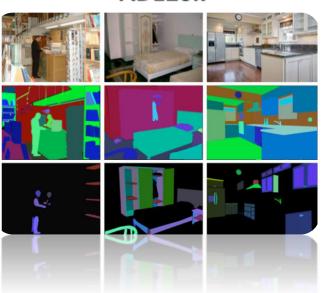


Image Classification on ImageNet-1K



- Settings:
 - Training on ImageNet-1K:
 - Optimizer: AdamW
 - Epochs: 300 with cosine decay learning rate
 - Warm-up: 20 epochs linear
 - **Batch size**: 1024
 - Initial learning rate: 0.001
 - Weight decay: 0.05
 - Pre-training on ImageNet-22K and fine-tuning on ImageNet-1K:
 - Optimizer: AdamW
 - **Pre-training**: 90 epochs, fine-tuning: 30 epochs
 - Batch size: Pre-training 4096, fine-tuning 1024
 - Learning rate: 0.001 for pre-training, 10^-5 for fine-tuning
 - Weight decay: 0.01 for pre-training, 10^-8 for fine-tuning

Image Classification on ImageNet-1K



Results:

- Comparisons with other backbones:
 - Swin-T vs. DeiT-S: +1.5 top-1 acc. (81.3 vs. 79.8 top-1 acc.) using 224x224 input
 - Swin-B vs. DeiT-B: +1.5 top-1 acc. /1.4 top-1 acc. (83.3/84.5 top-1 acc. vs. 81.8/83.1 top-1 acc.) using
 224x224/384x384 input
 - Swin-T vs. RegNetY-8G: -0.4 top-1 acc. (81.3 vs. 81.7 top-1 acc.) using 224x224 input
 - Swin-T vs. EffNet-B3: -0.3 top-1 acc. (81.3 vs. 81.6 top-1 acc.) using 224x224 input
- ImageNet-22K Pre-training Results:
 - **Swin-B:** 86.4 top-1 accuracy (+2.4 top-1 acc. higher than ViT-B/16 and +2 top-1 acc. higher than R-101x3)
 - Swin-L: 87.3 top-1 accuracy (+2.1 top-1 acc. better than ViT-L/16 and +2.9 top-1 acc. higher than R-101x3)

(a) Regular ImageNet-1K trained models									
method	image	#param.	EI ODa	throughput	ImageNet				
method	size	#param.	. FLOPS	(image / s)	top-1 acc.				
RegNetY-4G [48]	224 ²	21M	4.0G	1156.7	80.0				
RegNetY-8G [48]	224 ²	39M	8.0G	591.6	81.7				
RegNetY-16G [48]	224 ²	84M	16.0G	334.7	82.9				
EffNet-B3 [58]	300^{2}	12M	1.8G	732.1	81.6				
EffNet-B4 [58]	380^{2}	19M	4.2G	349.4	82.9				
EffNet-B5 [58]	456 ²	30M	9.9G	169.1	83.6				
EffNet-B6 [58]	528 ²	43M	19.0G	96.9	84.0				
EffNet-B7 [58]	600^{2}	66M	37.0G	55.1	84.3				
ViT-B/16 [20]	384 ²	86M	55.4G	85.9	77.9				
ViT-L/16 [20]	384 ²	307M	190.7G	27.3	76.5				
DeiT-S [63]	224 ²	22M	4.6G	940.4	79.8				
DeiT-B [63]	224 ²	86M	17.5G	292.3	81.8				
DeiT-B [63]	384 ²	86M	55.4G	85.9	83.1				
Swin-T	224 ²	29M	4.5G	755.2	81.3				
Swin-S	224 ²	50M	8.7G	436.9	83.0				
Swin-B	224 ²	88M	15.4G	278.1	83.5				
Swin-B	384 ²	88M	47.0G	84.7	84.5				

(b) ImageNet-22K pre-trained models

method	image	#param.	FI OPs	throughput	
method	size	"Paraiii.	LOIS	(image / s)	top-1 acc.
R-101x3 [38]	384^{2}	388M	204.6G	-	84.4
R-152x4 [38]	480^{2}	937M	840.5G	-	85.4
ViT-B/16 [20]	384 ²	86M	55.4G	85.9	84.0
ViT-L/16 [20]	384^{2}	307M	190.7G	27.3	85.2
Swin-B	224^{2}	88M	15.4G	278.1	85.2
Swin-B	384^{2}	88M	47.0G	84.7	86.4
Swin-L	384^{2}	197M	103.9G	42.1	87.3

Table 1. Comparison of different backbones on ImageNet-1K classification. Throughput is measured using the GitHub repository of [68] and a V100 GPU, following [63].

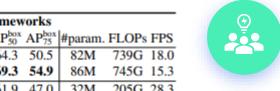
Object Detection on COCO



- Settings:
 - Models Evaluated:
 - Faster R-CNN
 - Mask R-CNN
 - Cascade Mask R-CNN
 - Training Configuration:
 - Backbone: Swin-T, Swin-B, Swin-L
 - Optimizer: AdamW
 - Learning Rate Schedule: Cosine decay
 - Batch Size: 16
 - **Epochs**: 36 (3x schedule)

Object Detection on COCO

- **Results:**
 - Swin-T vs. ResNet-50:
 - Achieves +3.4 to 4.2 box AP gains over ResNet-50.
 - **Comparison with Other Backbones:**
 - ResNet-50 (R-50):
 - Swin-T achieves +4.2 box AP over ResNet-50.
 - DeiT-S:
 - Swin-T achieves +2.5 box AP over DeiT-S.
 - ResNeXt101-32:
 - Swin-S shows +3.7 box AP over ResNeXt101-32.
 - ResNeXt101-64:
 - Swin-B shows +3.6 box AP over ResNeXt101-64.



	(a) Various frameworks									
Method	Backbone	AP ^{box}	AP ₅₀	AP ₇₅	#param.	FLOPs FPS				
Cascade	R-50	46.3	64.3	50.5	82M	739G 18.0				
Mask R-CNN	Swin-T	50.5	69.3	54.9	86M	745G 15.3				
ATSS	R-50	43.5	61.9	47.0	32M	205G 28.3				
Alss	Swin-T	47.2	66.5	51.3	36M	215G 22.3				
DonDointo V2	R-50	46.5	64.6	50.3	42M	274G 13.6				
RepPointsV2	Swin-T	50.0	68.5	54.2	45M	283G 12.0				
Sparse	R-50	44.5	63.4	48.2	106M	166G 21.0				
R-CNN	Swin-T	47.9	67.3	52.3	110M	172G 18.4				

(b) Various backbones w. Cascade Mask R-CNN APbox DeiT-S[†] | 48.0 67.2 51.7 | 41.4 64.2 44.3 | 80M 889G 10.4 R50 | 46.3 64.3 50.5 | 40.1 61.7 43.4 | 82M 739G 18.0 Swin-T 50.5 69.3 54.9 43.7 66.6 47.1 86M 745G 15.3 X101-32 48.1 66.5 52.4 41.6 63.9 45.2 101M 819G 12.8 Swin-S 51.8 70.4 56.3 44.7 67.9 48.5 107M 838G 12.0 X101-64 48.3 66.4 52.3 41.7 64.0 45.1 140M 972G 10.4 Swin-B 51.9 70.9 56.5 45.0 68.4 48.7 145M 982G 11.6

(c)	System-	level	Comp	parisoi

Method		ii-val AP ^{mask}		t-dev AP ^{mask}	#param.	FLOPs
RepPointsV2* [12]	-	-	52.1	-	-	-
GCNet* [7]	51.8	44.7	52.3	45.4	-	1041G
RelationNet++* [13]	-	-	52.7	-	-	-
SpineNet-190 [21]	52.6	-	52.8	-	164M	1885G
ResNeSt-200* [78]	52.5	-	53.3	47.1	-	-
EfficientDet-D7 [59]	54.4	-	55.1	-	77M	410G
DetectoRS* [46]	-	-	55.7	48.5	-	-
YOLOv4 P7* [4]	-	-	55.8	-	-	-
Copy-paste [26]	55.9	47.2	56.0	47.4	185M	1440G
X101-64 (HTC++)	52.3	46.0	-	-	155M	1033G
Swin-B (HTC++)	56.4	49.1	-	-	160M	1043G
Swin-L (HTC++)	57.1	49.5	57.7	50.2	284M	1470G
Swin-L (HTC++)*	58.0	50.4	58.7	51.1	284M	-

Table 2. Results on COCO object detection and instance segmentation. †denotes that additional decovolution layers are used to produce hierarchical feature maps. * indicates multi-scale testing.

Semantic Segmentation on ADE20K



Settings:

Dataset: ADE20K

Base Framework: UperNet

Results:

- Swin-S
 - Achieves 49.3 mIoU, which is +5.3 mIoU higher than DeiT-S (44.0) with similar computation cost.
 - Outperforms ResNet-101 by +4.4 mIoU.
 - Outperforms ResNeSt-101 by +2.4 mIoU.

Swin-L:

- With ImageNet-22K pre-training, achieves 53.5 mIoU on the validation set.
- Surpasses the previous best model (SETR) by <u>+3.2 mIoU</u>, which had 50.3 mIoU with a larger model size.

K	val	test			
ADE20K		test	#param.	FI OPs	FPS
Backbone	mIoU	score	трагані.	1 LOI 3	115
lesNet-101	45.2	-	69M	1119G	15.2
lesNet-101	44.1	-	63M	1021G	16.0
lesNet-101	45.9	38.5	-		
esNet-101	46.0	56.2	69M	1249G	14.8
esNet-101	45.3	56.0	56M	923G	19.3
esNet-101	44.9	-	86M	1029G	20.1
RNet-w48	45.7	-	71M	664G	12.5
esNeSt-101	46.9	55.1	66M	1051G	11.9
esNeSt-200	48.4	-	88M	1381G	8.1
T-Large [‡]	50.3	61.7	308M	-	-
DeiT-S [†]	44.0	-	52M	1099G	16.2
Swin-T	46.1	-	60M	945G	18.5
Swin-S	49.3	-	81M	1038G	15.2
Swin-B [‡]	51.6	-	121M	1841G	8.7
Swin-L [‡]	53.5	62.8	234M	3230G	6.2
֡	desNet-101 desNet-101 desNet-101 desNet-101 desNet-101 desNet-101 desNet-101 desNeSt-101 desNeSt-200 T-Large [‡] DeiT-S [†] Swin-T Swin-S Swin-B [‡]	tesNet-101 45.2 tesNet-101 44.1 tesNet-101 45.9 tesNet-101 46.0 tesNet-101 45.3 tesNet-101 44.9 IRNet-w48 45.7 tesNeSt-101 46.9 tesNeSt-200 48.4 T-Large [‡] 50.3 DeiT-S [†] 44.0 Swin-T 46.1 Swin-S 49.3 Swin-B [‡] 51.6	tesNet-101	Backbone	Ses

Table 3. Results of semantic segmentation on the ADE20K val and test set. † indicates additional deconvolution layers are used to produce hierarchical feature maps. ‡ indicates that the model is pre-trained on ImageNet-22K.

Ablation Studies



- Importance of Shifted Windows:
 - Comparison with single window partitioning shows:
 - +1.1% top-1 accuracy on ImageNet-1K
 - +2.8 box AP and +2.2 mask AP on COCO
 - <u>+2.8 mloU</u> on ADE20K
- Effect of Relative Position Bias:
 - Improves performance across all tasks:
 - <u>+1.2%/+0.8%</u> top-1 accuracy on ImageNet-1K
 - <u>+1.3/+1.5 box AP</u> and +1.1/+1.3 mask AP on COCO
 - <u>+2.3/+2.9 mloU</u> on ADE20K

	ImageNet		CC	CO	ADE20k
	top-1	top-5	APbox	AP^{mask}	mIoU
w/o shifting	80.2	95.1	47.7	41.5	43.3
shifted windows	81.3	95.6	50.5	43.7	46.1
no pos.	80.1	94.9	49.2	42.6	43.8
abs. pos.	80.5	95.2	49.0	42.4	43.2
abs.+rel. pos.	81.3	95.6	50.2	43.4	44.0
rel. pos. w/o app.	79.3	94.7	48.2	41.9	44.1
rel. pos.	81.3	95.6	50.5	43.7	46.1

Table 4. Ablation study on the *shifted windows* approach and different position embedding methods on three benchmarks, using the Swin-T architecture. w/o shifting: all self-attention modules adopt regular window partitioning, without *shifting*; abs. pos.: absolute position embedding term of ViT; rel. pos.: the default settings with an additional relative position bias term (see Eq. (4)); app.: the first scaled dot-product term in Eq. (4).





Conclusions

• Summary:

- Introduced a hierarchical vision Transformer using shifted windows.
- Addressed limitations of standard ViTs by incorporating hierarchical feature maps and shifted windows.
- Achieves linear computational complexity with respect to input image size.
- Outperforms previous methods in COCO object detection and ADE20K semantic segmentation.

Key Innovations:

- Shifted Window Mechanism: Effective and efficient for vision tasks.
- **Hierarchical Feature Maps:** Enables efficient computation and better handling of high-resolution images.
- Relative Position Bias: Enhances spatial understanding and performance.





Optimization: Further research into more efficient training methods and architectures.

Future Directions



Applications: Explore the use of shifted window-based self-attention in NLP.



Integration: Combining Swin Transformers with other models and techniques for enhanced performance.

End

Thanks for listening!